

```

import numpy as np

# Generate synthetic data
np.random.seed(42) # for reproducibility

# Number of samples and features
num_samples = 1000
num_features = 10

# Generate random feature vectors
X_train = np.random.randn(num_samples, num_features)

# Generate corresponding labels (binary classification)
y_train = np.random.randint(2, size=(num_samples, 1))

# Print shapes for verification
print("X_train shape:", X_train.shape)
print("y_train shape:", y_train.shape)

↩ X_train shape: (1000, 10)
  y_train shape: (1000, 1)

# Define neural network architecture
layer_sizes = [num_features, 2, 4, 8, 1] # Including input and output layers
num_layers = len(layer_sizes) - 1 # Excluding input layer

# Initialize parameters
parameters = {}
for l in range(1, num_layers + 1):
    parameters['W' + str(l)] = np.random.randn(layer_sizes[l], layer_sizes[l-1]) * 0.01
    parameters['b' + str(l)] = np.zeros((layer_sizes[l], 1))

# Sigmoid activation function
def sigmoid(z):
    return 1 / (1 + np.exp(-z))

# Forward propagation
def forward_propagation(X, parameters):
    cache = {'A0': X.T}

    for l in range(1, num_layers + 1):
        cache['Z' + str(l)] = np.dot(parameters['W' + str(l)], cache['A' + str(l-1)]) + parameters['b' + str(l)]
        cache['A' + str(l)] = sigmoid(cache['Z' + str(l)])

    return cache

# Binary cross-entropy loss
def binary_cross_entropy_loss(A, Y):
    m = Y.shape[0]
    loss = -1/m * np.sum(Y * np.log(A) + (1 - Y) * np.log(1 - A))
    return loss

# Backpropagation
def backward_propagation(cache, parameters, X, Y):
    m = Y.shape[0]
    grads = {}
    dZ = cache['A' + str(num_layers)] - Y.T

    for l in range(num_layers, 0, -1):
        grads['dW' + str(l)] = 1/m * np.dot(dZ, cache['A' + str(l-1)].T)
        grads['db' + str(l)] = 1/m * np.sum(dZ, axis=1, keepdims=True)
        dZ = np.dot(parameters['W' + str(l)].T, dZ) * cache['A' + str(l-1)] * (1 - cache['A' + str(l-1)])

    return grads

# Update parameters using gradient descent
def update_parameters(parameters, grads, learning_rate):
    for l in range(1, num_layers + 1):
        parameters['W' + str(l)] -= learning_rate * grads['dW' + str(l)]
        parameters['b' + str(l)] -= learning_rate * grads['db' + str(l)]

```

```

# Training hyperparameters
learning_rate = 0.01
num_epochs = 1000

# Training loop
for epoch in range(num_epochs):
    # Forward propagation
    cache = forward_propagation(X_train, parameters)

    # Calculate loss
    loss = binary_crossentropy_loss(cache['A' + str(num_layers)].T, y_train)

    # Backpropagation
    grads = backward_propagation(cache, parameters, X_train, y_train)

    # Update parameters
    update_parameters(parameters, grads, learning_rate)

    # Print loss every 100 epochs
    if epoch % 100 == 0:
        print(f"Epoch {epoch}, Cost: {loss}")

```

```

Epoch 0, Cost: 0.6931516932333045
Epoch 100, Cost: 0.6931341691822668
Epoch 200, Cost: 0.6931302859331474
Epoch 300, Cost: 0.6931294254084869
Epoch 400, Cost: 0.6931292347147695
Epoch 500, Cost: 0.6931291924564577
Epoch 600, Cost: 0.6931291830918574
Epoch 700, Cost: 0.6931291810166236
Epoch 800, Cost: 0.693129180556743
Epoch 900, Cost: 0.6931291804548314

```

```

# Training loop with outputs
for epoch in range(num_epochs):
    # Forward propagation
    cache = forward_propagation(X_train, parameters)

    # Calculate loss
    loss = binary_crossentropy_loss(cache['A' + str(num_layers)].T, y_train)

    # Backpropagation
    grads = backward_propagation(cache, parameters, X_train, y_train)

    # Update parameters
    update_parameters(parameters, grads, learning_rate)

    # Print outputs every 100 epochs
    if epoch % 100 == 0:
        print(f"Epoch {epoch}, Loss: {loss}")
        print("Forward Propagation Outputs:")
        for key, value in cache.items():
            print(key, ":", value)
        print("Backward Propagation Gradients:")
        for key, value in grads.items():
            print(key, ":", value)
        print("\n")

```

