

✓ Lab 11

Submitted by: Ahmed Khalid 200148

Submitted to: Sir Awais

✓ Task done in Lab

```
import random
import numpy as np

data = np.array([[4,21], [4,17], [14,24], [12,21], [5,19], [3,16], [6,22], [10,24], [11,25], [10,21]])
data = np.array(data)
K = 2

def distance(point1, point2):
    return np.sqrt(np.sum((point1 - point2) ** 2))

num_points = data.shape[0]

initial_indices = random.sample(range(num_points), K)
centroids = data[initial_indices]

def assign_clusters(data, centroids):
    clusters = [[] for _ in range(K)]
    for point in data:
        distances = [distance(point, centroid) for centroid in centroids]
        closest_centroid = np.argmin(distances)
        clusters[closest_centroid].append(point)
    return clusters

def update_centroids(clusters):
    new_centroids = []
    for cluster in clusters:
        new_centroid = np.mean(cluster, axis=0) if cluster else np.array([0,0])
        new_centroids.append(new_centroid)
    return np.array(new_centroids)

def kmeans(data, centroids, max_iters=100):
    for i in range(max_iters):
        clusters = assign_clusters(data, centroids)
        new_centroids = update_centroids(clusters)

        if np.all(centroids == new_centroids):
            print(f"Converged after {i+1} iterations")
            break

        centroids = new_centroids

    return centroids, clusters
```

```
final_centroids, final_clusters = kmeans(data, centroids)

print("Final centroids:\n", final_centroids.tolist())
for idx, cluster in enumerate(final_clusters):
    print(f"Cluster {idx+1}:\n", np.array(cluster).tolist())
```

```

Converged after 2 iterations
Final centroids:
[[11.4, 23.0], [4.4, 19.0]]
Cluster 1:
[[14, 24], [12, 21], [10, 24], [11, 25], [10, 21]]
Cluster 2:
[[4, 21], [4, 17], [5, 19], [3, 16], [6, 22]]

```

✓ Task done at home using penguins dataset

```
# Import Required Packages
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
import numpy as np
import random
```

```
# Load Dataset
penguins_df = pd.read_csv("penguins.csv")
```

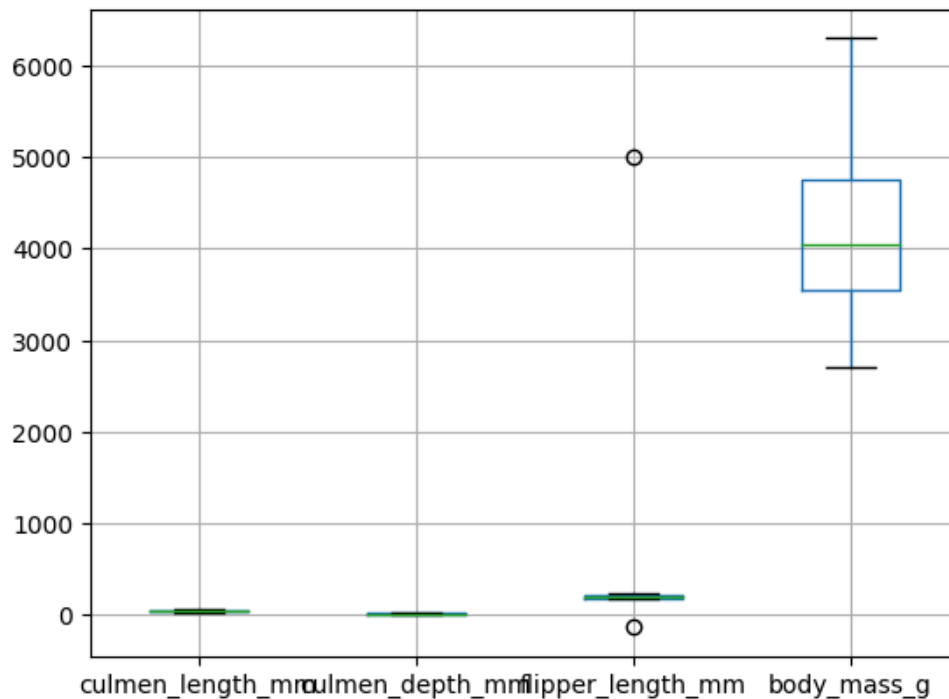
```
# Loading and examining the dataset
print(penguins_df.head())
```

```

culmen_length_mm  culmen_depth_mm  flipper_length_mm  body_mass_g  sex
0                39.1             18.7             181.0       3750.0  MALE
1                39.5             17.4             186.0       3800.0  FEMALE
2                40.3             18.0             195.0       3250.0  FEMALE
3                 NaN             NaN              NaN          NaN    NaN
4                36.7             19.3             193.0       3450.0  FEMALE

```

```
# Dealing with null values and outliers
penguins_df.boxplot()
plt.show()
```



```
penguins_df = penguins_df.dropna()
penguins_df[penguins_df["flipper_length_mm"] > 4000]
penguins_df[penguins_df["flipper_length_mm"] < 0]
penguins_clean = penguins_df.drop([9, 14])
```

```
# Perform preprocessing steps on the dataset to create dummy variables
df = pd.get_dummies(penguins_clean).drop("sex_", axis=1)
print(df.head())
```



	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	\
0	39.1	18.7	181.0	3750.0	
1	39.5	17.4	186.0	3800.0	
2	40.3	18.0	195.0	3250.0	
4	36.7	19.3	193.0	3450.0	
5	39.3	20.6	190.0	3650.0	

	sex_FEMALE	sex_MALE
0	False	True
1	True	False
2	True	False
4	True	False
5	False	True

```
# Perform preprocessing steps on the dataset - scaling
scaler = StandardScaler()
X = scaler.fit_transform(df)
penguins_preprocessed = pd.DataFrame(data=X, columns=df.columns)
print(penguins_preprocessed.head(10))
```



	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	\
0	-0.905520	0.793126	-1.428125	-0.569709	
1	-0.831938	0.128503	-1.071522	-0.507579	
2	-0.684775	0.435252	-0.429637	-1.191006	
3	-1.347011	1.099875	-0.572278	-0.942487	
4	-0.868729	1.764498	-0.786240	-0.693968	
5	-0.942311	0.333002	-1.428125	-0.725033	
6	-0.887125	1.253249	-0.429637	0.579691	
7	-0.537611	0.230753	-1.356804	-1.253136	
8	-0.997497	2.071247	-0.714919	-0.507579	
9	-1.365406	0.333002	-1.142843	-0.631839	

	sex_FEMALE	sex_MALE
0	-0.991031	0.997001
1	1.009050	-1.003008
2	1.009050	-1.003008
3	1.009050	-1.003008
4	-0.991031	0.997001
5	1.009050	-1.003008
6	-0.991031	0.997001
7	1.009050	-1.003008
8	-0.991031	0.997001
9	1.009050	-1.003008

```
# Perform PCA
pca = PCA(n_components=None)
dfx_pca = pca.fit(penguins_preprocessed)
print(dfx_pca.explained_variance_ratio_)
n_components = sum(dfx_pca.explained_variance_ratio_ > 0.1)
pca = PCA(n_components=n_components)
penguins_PCA = pca.fit_transform(penguins_preprocessed)
print(n_components)
```

2

```
# Number of clusters
K = 4
```

```
# Function to compute the Euclidean distance between two points
def distance(point1, point2):
    return np.sqrt(np.sum((point1 - point2) ** 2))
```

```
# Number of data points
num_points = penguins_PCA.shape[0]
```

```
# Randomly initialize centroids by selecting K data points
initial_indices = random.sample(range(num_points), K)
centroids = penguins_PCA[initial_indices]
```

```
# Function to assign points to the nearest centroid
def assign_clusters(data, centroids):
    clusters = [[] for _ in range(K)]
    for point in data:
        distances = [distance(point, centroid) for centroid in centroids]
        closest_centroid = np.argmin(distances)
        clusters[closest_centroid].append(point)
    return clusters
```

```
# Function to update centroids
def update_centroids(clusters):
    new_centroids = []
    for cluster in clusters:
        new_centroid = np.mean(cluster, axis=0) if cluster else np.array([0] * len(clusters[0]))
        new_centroids.append(new_centroid)
    return np.array(new_centroids)
```

```

# Function to perform K-means clustering
def kmeans(data, centroids, max_iters=100):
    for i in range(max_iters):
        clusters = assign_clusters(data, centroids)
        new_centroids = update_centroids(clusters)

        # Check for convergence
        if np.all(centroids == new_centroids):
            print(f"Converged after {i+1} iterations")
            break

        centroids = new_centroids

    return centroids, clusters

# Run K-means clustering
final_centroids, final_clusters = kmeans(penguins_PCA, centroids)

# Plotting the clusters
plt.scatter(penguins_PCA[:, 0], penguins_PCA[:, 1], c=KMeans(n_clusters=K).fit_predict(penguins_PCA), c
plt.scatter(final_centroids[:, 0], final_centroids[:, 1], marker='x', color='red', label='Centroids')
plt.xlabel("First Principal Component")
plt.ylabel("Second Principal Component")
plt.title(f"K-means Clustering (K={K})")
plt.legend()
plt.show()


# Print final centroids and clusters
print("\nFinal centroids:\n", final_centroids)
for idx, cluster in enumerate(final_clusters):
    print(f"Cluster {idx+1}:\n", np.array(cluster))

```



Converged after 3 iterations

/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default warnings.warn()



Final centroids:

```
[[-2.02559162  0.33725103]
 [ 0.29618153  2.27350811]
 [ 2.82671655  0.30596508]
 [ 0.25314985 -1.78195008]]
```

Cluster 1:

```
[[-1.9293044  0.4121497 ]
 [-1.95352479 0.3084854 ]
 [-2.28817761 -0.15827773]
 [-2.29152165 0.12644555]
 [-2.33416049 0.15078107]
 [-2.29097852 0.19172242]
 [-2.03556313 0.01409269]
 [-2.7868264  -0.14276205]
 [-2.80192041 -0.25463941]
 [-2.25960046 -0.14764826]
 [-2.35187589 -0.10802186]
 [-2.23509243 0.1561931 ]
 [-3.04481322 -0.44846709]
 [-2.49485913 0.32855442]
 [-2.20946064 0.215237 ]
 [-2.1372535  0.5712692 ]
 [-2.16582005 -0.07477592]
 [-2.72211005 -0.44061407]
 [-2.76587976 -0.06557282]
 [-2.74083453 -0.15451531]
 [-2.62934802 0.31844414]
 [-2.33397377 0.20191942]
 [-2.13533956 0.25063315]
 [-2.41393403 0.18461162]
 [-2.92081937 -0.07960675]
 [-2.13554517 0.31439783]
 [-2.87704097 0.28124726]
 [-2.63837943 0.33062618]
 [-2.20976996 0.43626328]
 [-2.82262097 0.1888527 ]
 [-2.12887198 0.81182706]
 [-2.49022642 0.51452888]
 [-2.53768100 0.15000530]]
```